

chapter 9 In-Class

Chapter 9 - Time series

```
library("Ecdat")
```

```
## Loading required package: Ecfun
##
## Attaching package: 'Ecfun'
## The following object is masked from 'package:base':
##
##      sign
##
## Attaching package: 'Ecdat'
## The following object is masked from 'package:datasets':
##
##      Orange
```

```
library("fGarch")
```

```
## Loading required package: timeDate
## Loading required package: timeSeries
## Loading required package: fBasics
##
## Rmetrics Package fBasics
## Analysing Markets and calculating Basic Statistics
## Copyright (C) 2005-2014 Rmetrics Association Zurich
## Educational Software for Financial Engineering and Computational Science
## Rmetrics is free software and comes with ABSOLUTELY NO WARRANTY.
## https://www.rmetrics.org --- Mail to: info@rmetrics.org
```

```
library("evir")
library("forecast")
```

```
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following object is masked from 'package:timeSeries':
##
##      time<-
##
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
## This is forecast 7.3
```

```
##
## Attaching package: 'forecast'

## The following object is masked from 'package:Ecfun':
##
##      BoxCox


data(CRSPday, package="Ecdat")
data(bmw)
summary(bmw)

##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## -0.1406000 -0.0066560  0.0000000  0.0003407  0.0071260  0.1172000

length(bmw)


## [1] 6146

plot(bmw)
```



```
acf(bmw)
```

Series bmw



The data seems to be stationary in the plot diagram. The autocorrelation function plot shows a well-behaved AR(1) or possibly AR(2) (here all changes lie within the dotted lines).

```
Box.test(x=bmw, lag=5, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data:  bmw
## X-squared = 44.987, df = 5, p-value = 1.46e-08
```

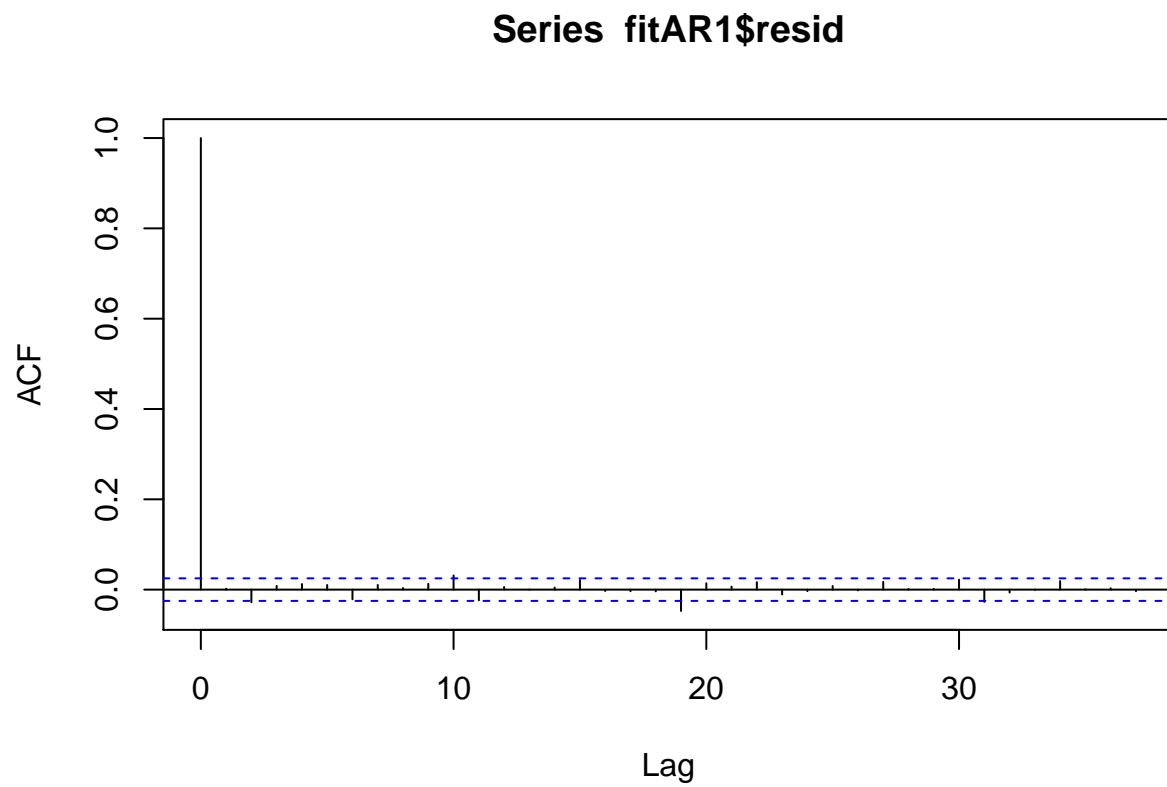
The test gives a very small p-value, so I could reject the null hypothesis that the data is independently distributed. There is serial correlation within the data.

So now


```
fitAR1 = arima(bmw, order = c(1,0,0))
# here I fit AR(1)
fitAR1
```

```
##
## Call:
## arima(x = bmw, order = c(1, 0, 0))
##
## Coefficients:
##          ar1  intercept
##       0.0811      3e-04
## s.e.  0.0127      2e-04
##
```

```
## sigma^2 estimated as 0.0002163: log likelihood = 17212.34, aic = -34418.68  
acf(fitAR1$resid)
```



```
plot(fitAR1$resid)
```



```
Box.test(fitAR1$resid,lag=5,type="Ljung-Box",fitdf=1)
```

```
##
## Box-Ljung test
##
## data: fitAR1$resid
## X-squared = 6.8669, df = 4, p-value = 0.1431
```

ACF graph goes quickly to zero, it should indicate stationarity. The Box test has a large p-value, which leads to fail to reject the null hypothesis that the residuals are uncorrelated, at least small lags. So, AR(1) provides a quite adequate fit here. Notice that we use AR(1)'s residuals to test to see if there is any autocorrelation left.

```
data(Mishkin)
summary(Mishkin)
```

```
##      pai1      pai3      tb1      tb3
## Min.   :-7.565  Min.   :-3.794  Min.    : 0.3215  Min.    : 0.5748
## 1st Qu.: 1.364  1st Qu.: 1.660  1st Qu.: 2.5395  1st Qu.: 2.8821
## Median : 3.589  Median : 3.679  Median : 4.5711  Median : 5.0470
## Mean   : 4.006  Mean    : 4.018  Mean    : 4.9983  Mean    : 5.4098
## 3rd Qu.: 6.118  3rd Qu.: 5.609  3rd Qu.: 6.8348  3rd Qu.: 7.4577
## Max.   :19.570  Max.    :16.431  Max.    :15.7906  Max.    :16.0278
##      cpi
## Min.   : 23.50
## 1st Qu.: 29.50
## Median : 39.00
## Mean   : 55.95
```

```
## 3rd Qu.: 84.40
## Max.    :133.80
```


```
?Mishkin
```

```
## starting httpd help server ...
```

```
## done
```


The first one is one-month inflation rate (in percent, annual rate).

```
inflation = Mishkin[,1]
plot(inflation)
```



```
# it doesn't seem stationary
acf(inflation)
```


Series inflation



The ACF decays to zero slowly, showing that this is a sign of either nonstationarity or possibly of s

```
fitAR1= arima(inflation, order=c(1,0,1))  
acf(fitAR1$resid)
```

Series fitAR1\$resid




```
Box.test(fitAR1$resid, lag=5, type="Ljung-Box", fitdf=1)
```

```
##  
## Box-Ljung test  
##  
## data: fitAR1$resid  
## X-squared = 22.491, df = 4, p-value = 0.00016
```

Low p-value show a significant autocorrelation within the past 4 lags.

```
library(timeSeries)  
df_inflation = diff(inflation)  
?diff  
acf(df_inflation) # quicky decaying to zero, quite stationary I assume
```



Series df_inflation



```
fitART=arima(df_inflation, order=c(1,0,0))  
fitAR1
```

```
##  
## Call:  
## arima(x = inflation, order = c(1, 0, 1))  
##  
## Coefficients:  
##          ar1          ma1  intercept  
##         0.9628      -0.7402         4.0090  
## s.e.   0.0182   0.0575         0.8936  
##  
## sigma^2 estimated as 8.741:  log likelihood = -1229.41,  aic = 2466.83  
acf(fitAR1$resid)
```

Series fitAR1\$resid



```
Box.test(fitAR1$residuals, lag=5, type="Ljung-Box", fitdf=1)
```

```
##
## Box-Ljung test
##
## data: fitAR1$residuals
## X-squared = 22.491, df = 4, p-value = 0.00016
```

there is some autocorrelation in the residuals for that small p-value as we can reject the null hypothesis


So now we are fitting AR(p) process to `df_inflation`. With `auto.arima`, R could find the optimal p, here `max.p = 20`

```
fit=auto.arima(df_inflation,max.p=20,max.q=0,max.d=0,max.P=0,max.D=0,max.Q=0,ic="aic")
fit
```

```
## Series: df_inflation
## ARIMA(8,0,0) with zero mean
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ar5      ar6      ar7
##      -0.6274  -0.4977  -0.5158  -0.4155  -0.3443  -0.2560  -0.1557
## s.e.   0.0456   0.0536   0.0576   0.0606   0.0610   0.0581   0.0543
##          ar8
##      -0.1051
## s.e.   0.0459
##
## sigma^2 estimated as 8.681: log likelihood=-1221.2
```

```
## AIC=2460.4   AICc=2460.78   BIC=2498.15
```

```
acf(fit$resid)
```



```
Box.test(fit$residuals, lag=5, type="Ljung-Box", fitdf=1)
```

```
##  
## Box-Ljung test  
##  
## data: fit$residuals  
## X-squared = 0.18324, df = 4, p-value = 0.9961
```

This fit seems to work so well that I cannot reject the non-autocorrelation test within the residuals. So `df_inflation` has AR(8) process.

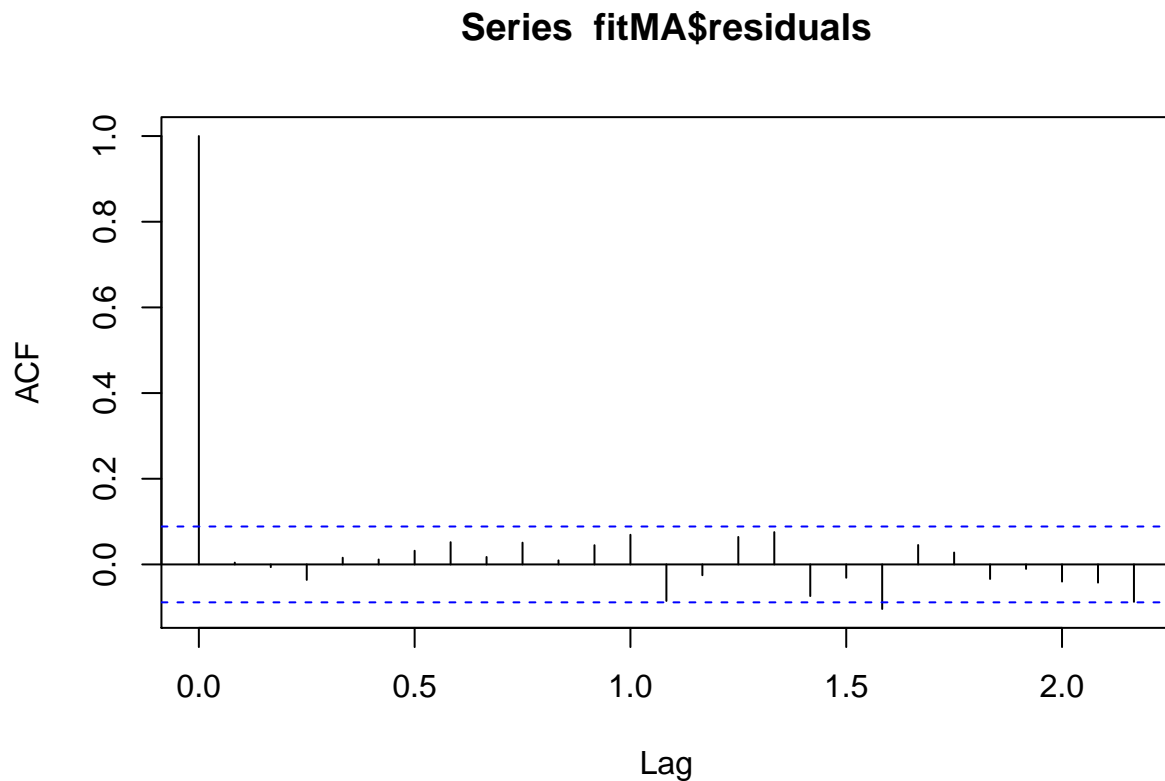
Next, we will fit MA (moving average process)

```
fitMA=arima(df_inflation, order = c(0,0,3))  
fitMA
```

```
##  
## Call:  
## arima(x = df_inflation, order = c(0, 0, 3))  
##  
## Coefficients:  
##          ma1          ma2          ma3  intercept  
##       -0.633   -0.1027   -0.1082   -0.0002  
## s.e.    0.046    0.0514    0.0470    0.0209  
##
```

```
## sigma^2 estimated as 8.505: log likelihood = -1220.26, aic = 2450.52
```

```
acf(fitMA$residuals)
```



```
Box.test(fitMA$residuals, lag=5, type="Ljung-Box", fitdf=1)
```

```
##
## Box-Ljung test
##
## data: fitMA$residuals
## X-squared = 0.86202, df = 4, p-value = 0.9299
```

Yup, the MA(3) fits well.

Now, we test with ARMA(p,q) fitting

```
fitARMA=auto.arima(inflation,max.p=5,max.q=5,max.d=0,max.P=0,max.Q=0,max.D=0,ic="aic")
fitARMA
```

```
## Series: inflation
## ARIMA(2,0,1) with non-zero mean
##
## Coefficients:
##          ar1      ar2      ma1  intercept
##          1.2074 -0.2237 -0.8573      4.1068
## s.e.    0.0587  0.0551  0.0337      1.0692
##
## sigma^2 estimated as 8.533: log likelihood=-1221.52
## AIC=2453.03 AICc=2453.16 BIC=2474.02
```

From above, we have $1.2074 > 1$, we might want to test for stationary process. Now we try to find root for $1 - ax - bx^2$ to see if the process is stationary. Note that if there exists a non-unit root, the process is stationary.

```
polyroot(c(1,-1.2074, 0.2237))
```

```
## [1] 1.021584+0i 4.375823-0i
```

Here we have two non-unit roots; thus, the process is stationary. We also have different tests for stationary process such as followings:

```
library("tseries")
adf.test(inflation)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: inflation
## Dickey-Fuller = -3.8651, Lag order = 7, p-value = 0.01576
## alternative hypothesis: stationary
```

```
pp.test(inflation)
```

```
## Warning in pp.test(inflation): p-value smaller than printed p-value
```

```
##
## Phillips-Perron Unit Root Test
##
## data: inflation
## Dickey-Fuller Z(alpha) = -248.75, Truncation lag parameter = 5,
## p-value = 0.01
## alternative hypothesis: stationary
```

```
kpss.test(inflation)
```

```
## Warning in kpss.test(inflation): p-value smaller than printed p-value
```

```
##
## KPSS Test for Level Stationarity
##
## data: inflation
## KPSS Level = 2.51, Truncation lag parameter = 5, p-value = 0.01
```

Just stationary like we stated.

Now we fit ARIMA process to the inflation data

```
fitARIMA=auto.arima(inflation,max.P=0,max.Q=0,max.D=0,ic="aic")
fitARIMA
```

```
## Series: inflation
## ARIMA(1,1,1)
##
## Coefficients:
##          ar1          ma1
##          0.2383 -0.8772
## s.e.    0.0550  0.0269
##
## sigma^2 estimated as 8.587: log likelihood=-1221.62
## AIC=2449.25 AICc=2449.29 BIC=2461.83
```


```
predict(fitARIMA,n.ahead=10)
```

```
## $pred
##           Jan           Feb           Mar           Apr           May           Jun           Jul
## 1991 3.706101 4.589298 4.799773 4.849930 4.861884 4.864732 4.865411
##           Aug           Sep           Oct
## 1991 4.865573 4.865611 4.865620
##
## $se
##           Jan           Feb           Mar           Apr           May           Jun           Jul
## 1991 2.930391 3.115623 3.175183 3.215216 3.250915 3.285350 3.319222
##           Aug           Sep           Oct
## 1991 3.352703 3.385842 3.418657
```


This process is parsimonious since we only need 3 parameters, reducing the problem of overfitting the model. We also predict the next 10 months based on this ARIMA(1,1,1) process. Notice that the predicted value is increasing with the also increasing standard deviation.

For this last part, we fit the inflation data with sample partial autocorrelation function.

```
pacf(bmw)
```



```
pacf(df_inflation)
```



The PACF is useful to identify the order of the AR process. The bmw log returns can be modeled as AR(1), while the df_inflation should try MA process instead, which is consistent with the analysis above.